SENG 365 Week 3 TypeScript and Data Persistence





The story so far

- What is a Web Application?
- HTTP
- JavaScript basics
- Asynchronous JavaScript
- Assignment 1



- TypeScript
- Data Persistence in Web Applications



TypeScript Handbook: https://www.typescriptlang.org/docs/handbook/intro.html



Problems with JavaScript

- Dynamically typed
- Type coercion behaves in unexpected ways
 - Poor IDE support
 - https://blog.campvanilla.com/javasc ript-the-curious-case-of-null-0-7b1 31644e274
- Different ECMAScript versions of the language that are not supported by all browsers

```
var container = "hello";
container = 43;
```

```
null > 0; // false
null == 0; // false
null >= 0; // true
```

TypeScript

- Developed by Microsoft
- Goal to create safer web code quicker
- Superset of JavaScript
 - All JavaScript code is TypeScript
- Adds:
 - Static typing
 - Type inference
 - Better IDE support
 - Strict null checking



TypeScript files

- TypeScript files use .ts extension
- Any JavaScript file can be converted to TypeScript by simply changing extension from .js to .ts
- The opposite is not true



Static typing in TypeScript

- Basic Types:
 - From JS primitives: boolean, number, bigint, string, array, tuple, object, null, undefined
 - Additional: enum, unknown, any, void, never

- Type declarations for variable
 - Do not change how the code runs
 - Are used by the compiler for type checking
 - Can be explicit or inferred by assignment

SP .

Static typing examples

```
let isDone: boolean = false;
let decimal: number = 6;
let hex: number = 0 \times f000d;
let binary: number = 0b1010;
let octal: number = 00744;
let big: bigint = 100n;
let color: string = "blue";
color = 'red';
```



Static typing examples

```
let list: number[] = [1, 2, 3];
let list: Array<number> = [1, 2, 3];
enum Color {
  Red,
  Green,
  Blue,
let c: Color = Color.Green;
```



Static typing functions

```
// Parameter type annotation
function greet(name: string) {
  console.log("Hello, " + name.toUpperCase() + "!!");
}

// Would be a runtime error if executed!
greet(42);

Argument of type 'number' is not assignable to parameter of type 'string'.
```

```
function getFavoriteNumber(): number {
  return 26;
}
```

void type is used when no return value



Static typing objects

- Duck typing based on the shape
- Can be anonymous or named using interface

```
function greet(person: { name: string; age: number }) {
  return "Hello " + person.name;
}
```

```
interface Person {
  name: string;
  age: number;
}

function greet(person: Person) {
  return "Hello " + person.name;
}
```

Properties can be optional using?

age?: number;



Static typing interfaces, types, and classes

Interface declarations can be used with classes

type is like interface but cannot be extended See https://cutt.ly/NAnFoG9

```
interface User {
 name: string;
  id: number;
class UserAccount {
 name: string;
  id: number;
  constructor(name: string, id: number) {
    this.name = name;
    this.id = id;
const user: User = new UserAccount("Murphy", 1);
```

Unions

```
function printId(id: number | string) {
  console.log("Your ID is: " + id);
}
// OK
printId(101);
// OK
printId("202");
// Error
printId({ myID: 22342 });
```

Union types

```
type WindowStates = "open" | "closed" | "minimized";
type LockStates = "locked" | "unlocked";
type PositiveOddNumbersUnderTen = 1 | 3 | 5 | 7 | 9;
```



Unions and typeof

```
function printId(id: number | string) {
 if (typeof id === "string") {
   // In this branch, id is of type 'string'
   console.log(id.toUpperCase());
 } else {
   // Here, id is of type 'number'
   console.log(id);
```

\$\frac{1}{2}

Strict null checking

```
let x: number = undefined;
```

Generates a compilation error

```
let x: number | undefined;
if (x !== undefined) x += 1; // this line will compile
x += 1; // this line will fail compilation
```

Compiling TypeScript

- Node.JS and browsers do not execute TypeScript
 - It must be compiled to JS first
- For Node.JS we need to add it to our project:
 - npm i -D typescript



TypeScript and Modules

- Node packages can have TypeScript bindings (supports IDE)
- Recall Node uses CommonJS modules (module.exports)
- Add .d.ts file to package

mymodule.ts

```
const maxInterval = 12;
function getArrayLength(arr) {
  return arr.length;
}
module.exports = {
  getArrayLength,
  maxInterval,
};
```

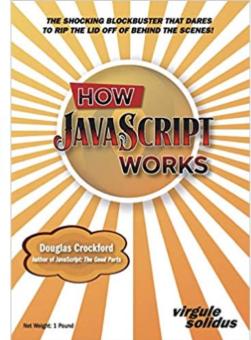
mymodule.d.ts

```
export function getArrayLength(arr: any[]): number;
export const maxInterval: 12;
```

Data in Web Applications

JSON, Relational DB, NoSQL





POJO and JSON

David Crockford's view: https://json.org/

A useful tool: https://json-to-js.com/ with npm version: npm i -g json-to-js

Useful tools (but not always accurate): https://tools.learningcontainer.com/

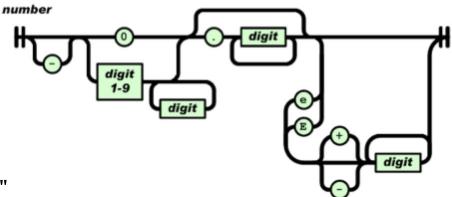


JSON Semi-formal definitions

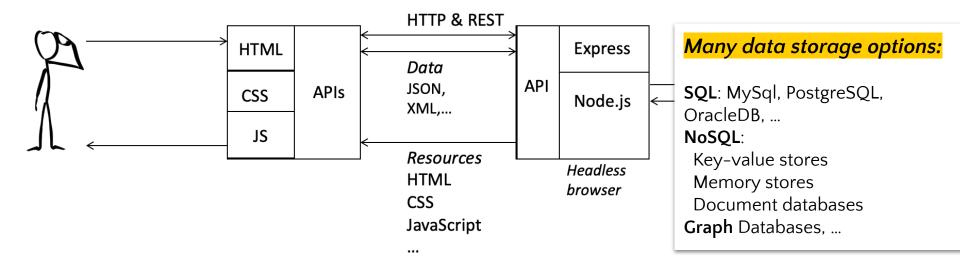
- JSON is a lightweight data-interchange format.
- A syntax for serializing data e.g., objects, arrays, numbers, strings, etc.
 - https://developer.mozilla.org/en-US/docs/Web/JavaScript /Reference/Global_Objects/JSON
- Data only, does not support comments except as a data field
- Not specific to JavaScript
 - Was originally intended for data interchange between Java and JavaScript
- No versioning for JSON: why?

JSON: some rules

- All key-names are double-quoted
- Values:
 - · Strings are double-quoted
 - Non-strings are not quoted
- Use \ to escape special characters, such as \ and "
- Numbers need to be handled carefully
 - · e.g. a decimal must have a trailing digit
 - Correct: 27.0
 - Incorrect 27.
 - Correct 27
- Can't shouldn't JSONify functions or methods
- See the following for guidance:
 - https://developer.mozilla.org/en-US/docs/Web/JavaScript/Reference/Global_Objects/JSON
 - https://json.org/
- And this for an... interesting discussion... on JSON syntax:
 - https://stackoverflow.com/questions/19176024/how-to-escape-special-characters-in-building-a-json-string



https://json.org/



Reference model

Machine

HTTP client

User

Human

X

HTTP Server

Machine

Database

Machine

Relational databases

- One of the few cases where a theoretical contribution in academic computer science led innovation in industry
- Relational model (E.F. Codd 1970)
 - Data is presented as relations
 - Collections of tables with columns and rows (tuples)
 - Each **tuple** has the same attributes
 - Unique key per row
 - **Relational algebra** that defines operations: UNION, INTERSECT, SELECT, JOIN, etc.

ACID database transactions

- Atomicity—"all or nothing" if one part of a transaction fails, then the whole transaction fails
- Consistency—the database is kept in a consistent state before and after transaction execution

 Isolation—one transaction should not see the effects of other, in progress, transactions

Durability—ensures transactions, once committed, are persistent

"The end of an architectural era?"

- Traditional RDBMSs
 - ACID properties were requirement for data handling

- Over the past few decades:
 - Moore's Law—CPU architectures have changed how they acquire speed
 - New requirements for data processing have emerged
 - Stonebraker et al. (2007), suggest that "one size fits all" DBs not sufficient
 - Still... relational databases are extremely useful in many cases

CAP Theorem

- In distributed computing, choose two of:
 - Consistency—every read receives the most recent data
 - Availability—every read receives a response
 - Partition tolerance—system continues if network goes down
- Situation is actually more subtle than implied above
 - Can adaptively choose appropriate trade-offs
 - Can understand semantics of data to choose safe operations

BASE

- Give up consistency (first part of CAP) and we can instead get:
 - Basic Availability—through replication
 - Soft state—the state of the system may change over time
 - This is due to the eventual consistency...
 - Eventual consistency—the data will be consistent eventually
 - ... if we wait long enough
 - (and probably only if data is not being changed frequently)

ACID versus **BASE** example (1/2)

- Suppose we wanted to track people's bank accounts:
 CREATE TABLE user (uid, name, amt_sold, amt_bought)
 CREATE TABLE transaction (tid, seller_id, buyer_id, amount)
- OACID transactions might look something like this:
 BEGIN
 INSERT INTO transaction(tid, seller_id, buyer_id, amount);
 UPDATE user SET amt_sold=amt_sold + amount WHERE
 id=seller_id;
 UPDATE user SET amt_bought=amt_bought + amount WHERE
 id=buyer_id;

END

ACID versus **BASE** Example (2/2)

If we consider amt_sold and amt_bought as estimates, transaction can be split: BFGTN INSERT INTO transaction(tid, seller_id, buyer_id, amount): FND **BFGTN** UPDATE user SET amt_sold=amt_sold + amount WHERE id=seller_id: UPDATE user SET amt_bought=amt_bought + amount WHERE id=buyer_id; **END**

- Consistency between tables is no longer guaranteed
- Failure between transactions may leave DB inconsistent



Key value databases overview

- Unstructured data (i.e., schema-less)
- Primary key is the only storage lookup mechanism
- No aggregates, no filter operations
- Simple operations such as:
 - Create—store a new key-value pair
 - Read—find a value for a given key
 - Update—change the value for a given key
 - Delete—remove the key-value pair



Key value databases

Advantages

- Simple
- Fast
- Flexible (able to store any serialisable data type)
- High scalability
- Can engineer high availability

Disadvantages

- Stored data is not validated
 - NOT NULL checks
 - colour versus color
- Complex to handle consistency
- Checking consistency becomes the application's problem
- No relationships—each value independent of all others
- No aggregates (SUM, COUNT, etc.)
- No searching (e.g., SQL SELECT-style) other than via key



Key value database implementations

- Amazon Dynamo (now DynamoDB)
- Oracle NoSQL Database, ... (eventually consistent)
- Berkeley DB, ... (ordered)
- Memcache, Redis, ... (RAM)
- LMDB (used by OpenLDAP, Postfix, InfluxDB)
- LevelDB (solid-state drive or rotating disk)
- IndexedDB (in the browser)



Dynamo Amazon's Highly Available Key-value Store

- Just two operations:
 - o put(key, context, object)
 - o get(key) → context, object
- Context provides a connection to DynamoDB
 - contains information not visible to caller
 - but is used internally, e.g., for managing versions of the object
- Objects are typically around 1MiB in size

<mark>Dynamo</mark> Design

- Reliability is one of the most important requirements
 - Significant financial consequences in its production use
 - Impacts user confidence
- Service Level Agreements (SLAs) are established
- Used within Amazon for:
 - best seller lists; shopping carts; customer preferences; session management; sales rank; product catalog



Redis in memory store

- Whole database is stored in RAM
 - Very fast access
 - Useful for cached data on the server
 - E.g. commonly accessed data from RDBMS can be stored in memory store on same computer as the API server.
- Key-value store where the value can be a complex data structure
 - Strings, Bitarrays, Lists, Sets, Hashes
 - Streams (useful for logs)
 - Binary-safe keys
 - Command set for optimized load, storing, and changing data values 36



Document databases

- Semi-structured data model
- Storage of documents:
- typically JSON or XML
- could be binary (PDF, DOC, XLS, etc.)
- Additional metadata (providence, security, etc.)
- Builds index from contexts and metadata



Document databases

Advantages

- Storage of raw program types (JSON/XML)
- Indexed by content and metadata
- Complex data can be stored easily
- No need for costly schema migrations
- (Always remember that your DB is likely to need to evolve!)

Disadvantages

- Same data replicated in each document
- Risk inconsistent or obsolete document structures



Document database implementations

- ElasticSearch
- LinkedIn's Espresso
- CouchDB
- MongoDB
- Solr / Apache Lucene
- RethinkDB
- Microsoft DocumentDB
- PostgreSQL (when used atypically)



- Node (or vertex)—represents an entity
- Edge—represents relationship between nodes
- Bidirectional (usually illustrated without arrowheads)
- Unidirectional (usually illustrated with an arrowhead)
- Properties—describe attributes of the node or edge
- Often stored as a key-value set
- Hypergraph one edge can join multiple nodes


Street map connectivity is a graph

Node


Traffic junction

• Edge

- Shows traffic flow
- Can be uni/bidirectional



Edges can have properties



name: "Union Street East"
type: "residential"

max speed: "50"

restrictions: "commercial"


surface: "tarmac"
status: "closed"
reason: "repairs"
length: "250m"

name: "Forth Street"
type: "residential"
max_speed: "50"
furniture: "bus stop"

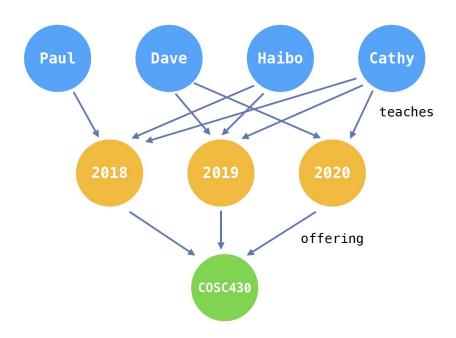
name: "Leithbank"
type: "residential"
max_speed: "50"
direction: "one way"

traffic_control: "traffic lights"

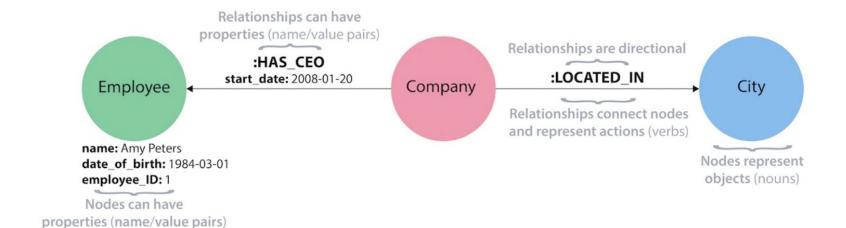
Nodes can have properties



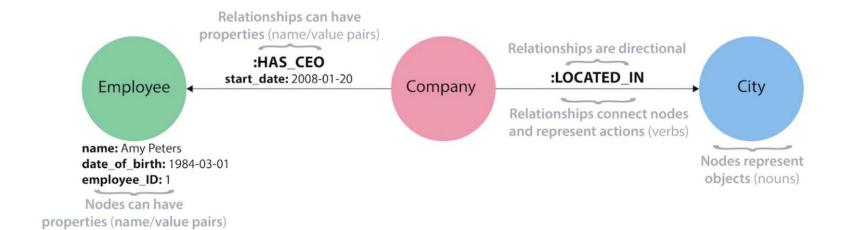
Different types of nodes

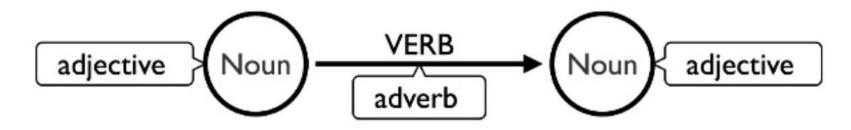


Building blocks of property graph model



Building blocks of property graph model





Why do we need graph databases?

- We can store graphs in RDBMSs, e.g.,
 - Node table
 - Edge table

- But, joins between nodes and edges are common
 - ... as the number of hops in a graph increases, this becomes increasingly expensive
- Some problems best suit direct representation in graphs
 - E.g. social graph



Designing graph databases

- Typical mapping from application's data to a graph:
 - Entities are represented as nodes
 - Connections are represented as edges between nodes
 - Connection semantics dictate directions of edges
 - Entity attributes become node properties
 - Link strength / weight / quality maps to relationship properties
- Other metadata will also be include in property sets
 - o e.g., information about data entry and revision



Graph database implementations

- Neo4j
 - https://neo4j.com/developer/graph-database/
- Amazon Neptune
- JanusGraph (scalable, distributed graph database)
- ArangoDB
- OrientDB
- RedisGraph (in memory)
- RDF-specific
 - Virtuoso, BlazeGraph, AllegroGraph
- Others... see https://tinkerpop.apache.org

Graph DB Query languages

- Cypher developed for neo4j but used by other systems
 - Declarative language (like SQL for graph databases)
 - Standard CRUD Create, Read, Update, Delete operations on the elements of the graph
 - Match patterns in the graph

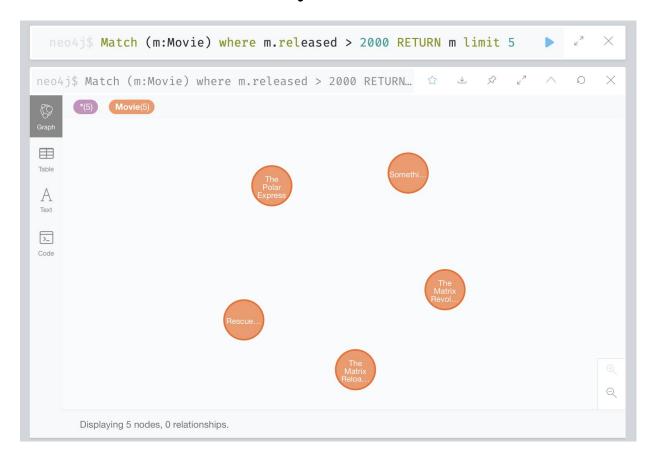
```
(node)-[:RELATIONSHIP]->(node)
```

```
(node {key: value})-[:RELATIONSHIP]->(node)
```

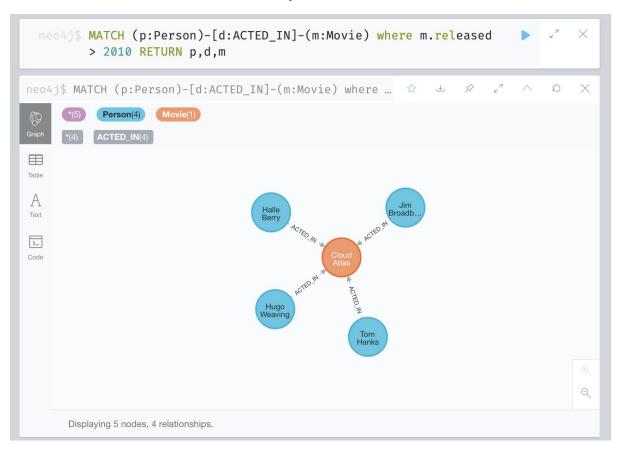
• Alternatives:

- SPARQL querying RDF graphs
- Gremlin graph traversal language for Apache Tinkerpop
- PGQL Oracle mix of SQL SELECT–style with graph matching

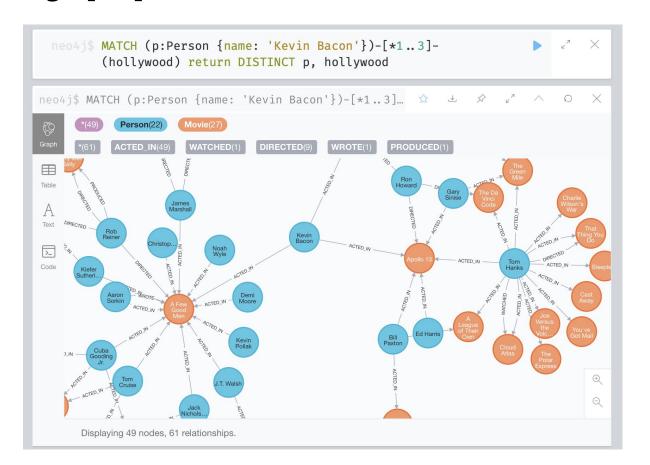
Cypher MATCH and RETURN keywords



Cypher MATCH and RETURN keywords



Complex graph queries



Connecting to neo4j from nodeJS

```
Shell

npm install neo4j-driver
```

```
JavaScript
                                                                                                    Copy to Clipboard
const neo4j = require('neo4j-driver')
const driver = neo4j.driver(uri, neo4j.auth.basic(user, password))
const session = driver.session()
const personName = 'Alice'
try {
  const result = await session.run(
    'CREATE (a:Person {name: $name}) RETURN a',
    { name: personName }
  const singleRecord = result.records[0]
  const node = singleRecord.get(0)
  console.log(node.properties.name)
} finally {
  await session.close()
// on application exit:
await driver.close()
```